

SIMATS ENGINEERING
ASSESSMENT TOOL 3

COURSE NAME: Database Management Systems

COURSE CODE: CSA05

COURSE OUTCOME COVERED

CO4: Design database solutions incorporating file organization, indexing, and hashing techniques to optimize data storage and retrieval efficiency. (BL3, BL4)

Activity Tool : Comparative Analysis (Low)

Assessment Tool Weightage: 10%

Code of Conduct:

I P SHASHANK (192524282) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.
I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the student:

P SHASHANK

Name: P SHASHANK **Reg. No:** 192524282

QUESTIONS

Q.No	Question	Bloom's Level
1	Compare sequential file organization and heap file organization in terms of data insertion, deletion, and retrieval efficiency.	4 – Analyze
2	Differentiate between clustered indexing and non-clustered indexing with suitable database examples.	4 – Analyze
3	Compare primary indexing and secondary indexing based on storage requirements and search performance.	4 – Analyze
4	Analyze the differences between static hashing and dynamic hashing in database systems.	4 – Analyze
5	Compare B-Tree indexing and B+ Tree indexing for large-scale database applications.	4 – Analyze
6	Differentiate linear hashing and extendible hashing in terms of bucket management and scalability.	4 – Analyze
7	Compare dense indexing and sparse indexing with respect to memory usage and access speed.	4 – Analyze
8	Analyze the advantages and limitations of indexed sequential file organization versus direct file organization.	4 – Analyze
9	Compare single-level indexing and multi-level indexing for efficient database retrieval operations.	4 – Analyze
10	Evaluate the performance of hashing techniques and indexing techniques for exact-match and range queries in databases.	5 – Evaluate

RUBRICS FOR CALCULATION:

Criteria	Weightage (%)	Level 4 (Exemplary)	Level 3 (Proficient)	Level 2 (Developing)	Level 1 (Beginning)
Scope of Items	20%	Selects highly relevant items with clear scope and justification	Selects appropriate items with minor gaps	Selection is limited or partially relevant	Items are unclear or not relevant
Comparison Depth	25%	Provides detailed and comprehensive comparison across multiple aspects	Provides adequate comparison with some depth	Limited comparison with few aspects	Very minimal or no comparison
Use of Evidence	20%	Supports comparison with accurate data, examples, or references	Uses some supporting evidence with minor gaps	Limited or weak supporting evidence	No supporting evidence provided
Critical Thinking	20%	Demonstrates strong critical thinking with clear	Shows reasonable interpretation with minor gaps	Basic interpretation; lacks depth	No meaningful interpretation

Criteria	Weightage (%)	Level 4 (Exemplary)	Level 3 (Proficient)	Level 2 (Developing)	Level 1 (Beginning)
		interpretation and reasoning			
Clarity of Presentation	15%	Well-organized, clear, and logically structured presentation	Generally clear with minor issues	Somewhat unclear or poorly structured	Disorganized and difficult to understand

Q1. Sequential File Organization vs Heap (Pile) File Organization

Bloom's Level: 4 – Analyze

Comparison Table

Aspect	Sequential File Organization	Heap File Organization
Definition	Records stored physically in order of a key field	Records stored in the order they arrive, no fixed order
Insertion	Slow — may require shifting/rewriting to keep order, or use of overflow area	Very fast — new record simply appended at the end
Deletion	Requires reorganization or tombstone/marker to keep file sorted	Fast — mark as deleted; space can be reused later
Retrieval (search)	Efficient for key-based search using binary search — $O(\log n)$	Inefficient — full linear scan needed — $O(n)$
Range queries	Very efficient since data is already sorted	Poor — every record must be scanned
Storage overhead	Low, but periodic reorganization needed	Minimal overhead but wasted space from deleted records
Best use case	Batch processing, reporting, payroll systems	Logging, staging/temporary data, bulk loading before indexing

Explanation

Sequential file organization stores records physically sorted by a key attribute (e.g., roll number), which allows fast key-based and range retrieval using binary search, but insertions and deletions are costly because maintaining sort order may require shifting records or using an overflow file that is periodically merged back.

Heap (pile) file organization simply appends new records wherever free space is available, with no ordering. This makes insertion an $O(1)$ operation, but every retrieval — even for a single record — typically requires scanning the entire file ($O(n)$), which becomes very expensive as the file grows.

Example(s)

- A student database sorted by roll number (sequential) supports fast rank-list generation, while a temporary log file that only needs to record events as they occur (heap) benefits from fast appends.

Justification / Critical Analysis

The choice between the two is a direct trade-off between write performance and read performance. Sequential organization is preferable when the workload is read-heavy and queries are key-based or range-based, whereas heap organization is preferable when the workload is write-heavy and records are rarely searched individually. In practice, heap files are often used as a staging area before bulk-loading into an indexed structure that combines fast writes with fast reads.

Q2. Clustered Indexing vs Non-Clustered Indexing

Bloom's Level: 4 – Analyze

Comparison Table

Aspect	Clustered Indexing	Non-Clustered Indexing
Physical order	Data rows are physically sorted and stored in the same order as the index	Index is a separate structure; data rows are not physically reordered
Number per table	Only one clustered index per table (since data can have only one physical order)	Multiple non-clustered indexes can exist on the same table
Storage	No separate data copy — index IS the data arrangement	Extra storage needed for index structure with pointers to actual rows
Retrieval speed	Faster for range queries because related rows are physically adjacent	Slightly slower — requires an extra lookup (pointer to row) after finding the index entry
Example (SQL Server)	Primary Key by default becomes a clustered index	CREATE INDEX on a non-key column, e.g., dept or email

Explanation

A clustered index determines the physical storage order of table rows — the leaf level of the index IS the data itself. Because of this, a table can have only one clustered index (commonly built on the primary key).

A non-clustered index is a separate structure containing the indexed column values along with a pointer (row locator) to where the actual row is physically stored. This allows several non-clustered indexes to coexist on the same table.

Example(s)

- In a Students table, roll_no (primary key) is typically the clustered index, so records are physically stored roll-number order — ideal for range reports like 'roll numbers 100 to 200'.
- A non-clustered index on the dept column allows fast lookup of 'all students in CSE' without reordering the whole table.

Justification / Critical Analysis

Below is real evidence generated using SQLite (a relational engine) on a 2,000-row Students table, showing how the query optimizer chooses different access paths depending on which column is indexed.

Code & Output Evidence (SQLite demonstration)

The following script creates a Students table (roll_no INTEGER PRIMARY KEY, name, dept, marks) with 2,000 rows and inspects the query execution plan before and after indexing dept.

```
CREATE TABLE students (roll_no INTEGER PRIMARY KEY, name TEXT, dept TEXT, marks INTEGER);
-- 2000 rows inserted

EXPLAIN QUERY PLAN SELECT * FROM students WHERE dept='CSE';
>> SCAN students                                -- (before creating index: full scan)

CREATE INDEX idx_dept ON students(dept);
EXPLAIN QUERY PLAN SELECT * FROM students WHERE dept='CSE';
>> SEARCH students USING INDEX idx_dept (dept=?)  -- non-clustered index used

EXPLAIN QUERY PLAN SELECT * FROM students WHERE roll_no=500;
>> SEARCH students USING INTEGER PRIMARY KEY (rowid=?)  -- clustered access via PK

Full table scan time (no index) : 0.195 ms
Indexed lookup time             : 0.056 ms    (~3.5x faster)
```

Output interpretation: Without an index, SQLite's query planner performs a full SCAN of every row. After CREATE INDEX idx_dept, the same query switches to a SEARCH using the index — a non-clustered, secondary structure. Lookup by roll_no (the PRIMARY KEY, which is clustered in SQLite's rowid table) is already an indexed SEARCH by default. The measured timing confirms indexed access is faster than a full scan even on a modest dataset, and the gap widens sharply as table size grows.

Q3. Primary Indexing vs Secondary Indexing

Bloom's Level: 4 – Analyze

Comparison Table

Aspect	Primary Indexing	Secondary Indexing
Built on	Ordering (primary/key) attribute of a sequentially organized file	Any non-ordering attribute, which may or may not be a candidate key
Type	Usually sparse — one index entry per data block	Usually dense — one index entry per record
Storage requirement	Lower, since fewer entries are stored (one per block)	Higher, since an entry is stored for every record
Search performance	Fast — index is small enough to often fit in memory; binary search on index then scan the block	Slower per index lookup due to larger index size, but enables search on non-key attributes
Data file dependency	Requires the data file to be physically sorted on the indexed field	Data file need not be sorted on the indexed field

Explanation

A primary index is built on the ordering key of a file that is itself physically sorted on that key. Because the data file is sorted, the index needs only one entry per block (sparse index), which keeps the index small and fast to search.

A secondary index is built on a field that is not used to physically order the file. Since the underlying data is not sorted by this field, the index must generally be dense — containing one entry for every record — so it can point directly to each matching record.

Example(s)

- roll_no in a Students table sorted by roll_no uses a primary (sparse) index.
- An index on email or phone_number in the same table, where rows are not physically sorted by those fields, must be a secondary (dense) index.

Justification / Critical Analysis

Primary indexes offer better storage efficiency and search speed but are restricted to the sorting attribute of the file, of which there can only be one. Secondary indexes trade extra storage for flexibility, letting applications efficiently query on any commonly searched attribute (e.g., searching by email even though the file is sorted by roll number).

Q4. Static Hashing vs Dynamic Hashing

Bloom's Level: 4 – Analyze

Comparison Table

Aspect	Static Hashing	Dynamic Hashing
Number of buckets	Fixed at file creation and does not change	Grows or shrinks automatically as data volume changes
Hash function	Fixed function producing a fixed range of bucket addresses	Extensible function (often uses extra bits of the hash value) that adapts to bucket growth
Handling growth	Bucket overflow chains form when data exceeds capacity, degrading performance	New buckets are added on demand, avoiding long overflow chains
Reorganization cost	Full file reorganization needed if resized (expensive)	Only local restructuring (e.g., splitting one bucket) needed
Performance over time	Degrades as overflow chains grow	Stays roughly constant since bucket count adapts to load
Complexity	Simple to implement	More complex — needs directory structures (as in extendible hashing)

Explanation

Static hashing uses a hash function that maps a fixed number of buckets, decided once when the file is created. If the number of records grows beyond the original capacity, records collide and are placed in overflow chains, which gradually slows down retrieval.

Dynamic hashing addresses this limitation by allowing the number of buckets to grow or shrink as the data volume changes, typically by using an expandable directory and progressively longer hash-value prefixes to route records to buckets.

Example(s)

- A static-hash table sized for 1,000 employee records will start performing poorly with long overflow chains once the company grows to 5,000 employees.
- A dynamic (extendible) hash index automatically splits the overloaded bucket and doubles the directory as needed, keeping lookups close to $O(1)$ even as the dataset grows.

Justification / Critical Analysis

Static hashing is simpler and has lower overhead, making it suitable for applications with a known, stable record count. Dynamic hashing is essential for databases with unpredictable or rapidly growing data, since it maintains near-constant-time access without requiring a full file reorganization.

Q5. B-Tree Indexing vs B+ Tree Indexing

Bloom's Level: 4 – Analyze

Comparison Table

Aspect	B-Tree Indexing	B+ Tree Indexing
Data storage location	Both keys and data (or record pointers) are stored in internal AND leaf nodes	Data pointers exist only in leaf nodes; internal nodes store only keys for routing
Leaf node linking	Leaf nodes are not linked to each other	Leaf nodes are linked in a sequence (linked list), enabling efficient range scans
Search for a single record	Can be faster since data may be found at an internal node before reaching a leaf	Always traverses to a leaf node, giving consistent, predictable search time
Range queries	Less efficient — no direct link between leaves, requires re-traversal	Very efficient — simply follow leaf-level linked list
Storage efficiency	Lower fan-out since internal nodes also carry data	Higher fan-out since internal nodes carry only keys, so trees are shallower
Common usage	Less common in modern RDBMS	Widely used — default index structure in MySQL (InnoDB), PostgreSQL, Oracle

Explanation

A B-Tree stores both keys and associated data pointers at every level of the tree (internal and leaf), which can make single-record lookups slightly faster if the target happens to reside near the root, but it complicates sequential/range access since data is scattered across levels.

A B+ Tree restricts actual data pointers to leaf nodes only; internal nodes act purely as a routing index. Critically, all leaf nodes are linked together, so once a range query locates its starting point, it can scan sequentially forward without re-traversing the tree — a major advantage for large-scale database applications.

Example(s)

- MySQL's InnoDB storage engine and PostgreSQL both implement B+ Trees for primary and secondary indexes precisely because most database queries (BETWEEN, ORDER BY, >, <) benefit from fast range scans.

Justification / Critical Analysis

For large-scale applications with frequent range queries, sorting, and reporting, B+ Trees are almost always preferred because of their linked-leaf structure and greater fan-out (shallower tree, fewer disk I/Os). B-Trees may still be acceptable when workloads are dominated by single-key exact lookups and range queries are rare, but in practice virtually all major relational database systems adopt B+ Trees as the default index structure.

Q6. Linear Hashing vs Extendible Hashing

Bloom's Level: 4 – Analyze

Comparison Table

Aspect	Linear Hashing	Extendible Hashing
Directory structure	No separate directory — buckets are split in a predetermined linear order	Uses a directory of pointers to buckets, indexed by bit-prefixes of the hash

Aspect	Linear Hashing	Extendible Hashing
Bucket splitting order	Buckets split in round-robin sequence regardless of which bucket overflowed	Only the specific overflowing bucket is split
Overflow handling	Uses overflow chaining until its turn to split arrives	Directory doubling (if needed) immediately accommodates the split
Directory growth	Not applicable — no directory to grow	Directory may double in size when local depth exceeds global depth
Memory overhead	Lower — no directory to maintain	Higher — must store and potentially resize the directory
Predictability of splits	Predictable, systematic split order	Split location is data-driven (based on which bucket is full)

Explanation

Linear hashing avoids the overhead of maintaining a directory by splitting buckets in a fixed, predetermined round-robin order — one bucket at a time — regardless of which bucket actually overflowed, relying on temporary overflow chains in the interim.

Extendible hashing maintains an explicit directory of pointers to buckets, addressed using the leading bits of a record's hash value. When a specific bucket overflows, only that bucket is split; if its local depth exceeds the directory's global depth, the directory itself doubles in size to accommodate finer addressing.

Example(s)

- Linear hashing is well suited to file systems that want predictable, low-overhead growth without an auxiliary directory.
- Extendible hashing is used in database index structures where quick, targeted splitting of only the overloaded bucket is valuable, at the cost of directory memory.

Justification / Critical Analysis

Both techniques aim to support dynamic growth without full file reorganization, but they trade off differently: linear hashing keeps memory overhead low with a simple, systematic splitting scheme but may temporarily tolerate overflow chains; extendible hashing reacts immediately and precisely to the bucket that actually overflowed but requires extra memory for the directory, which can double in size during heavy growth.

Q7. Dense Indexing vs Sparse Indexing

Bloom's Level: 4 – Analyze

Comparison Table

Aspect	Dense Indexing	Sparse Indexing
Index entries	One index entry for every record in the data file	One index entry per data block (not per record)
Memory usage	Higher — index size grows proportionally with the number of records	Lower — index size grows only with the number of blocks

Aspect	Dense Indexing	Sparse Indexing
Access speed (point query)	Faster — index directly points to the exact record	Slightly slower — index points to a block, then a scan within the block is needed
Applicability	Can be built on sorted or unsorted files	Requires the data file to be sorted on the indexed field
Storage trade-off	Uses more disk space for the index structure	More space-efficient, especially for large files

Explanation

A dense index contains an index record for every single search-key value in the data file, so it can directly locate any record without scanning, but this means the index grows as large as the number of records.

A sparse index contains an entry only for some records — typically one per disk block — so it is much smaller and can often be held in memory, but locating a specific record additionally requires a linear scan within the located block.

Example(s)

- A dense index on a Customers table's customer_id allows O(1)-style direct lookup of any customer, at the cost of index size equal to the table size.
- A sparse index on a sorted Orders table (one entry per block) saves substantial space when the table has millions of rows spread over relatively few blocks.

Justification / Critical Analysis

The decision is essentially a space-versus-directness trade-off. Dense indexing favors raw lookup speed at the cost of memory usage, making it attractive for smaller tables or when the file is unsorted. Sparse indexing favors compactness and is only viable on a file already sorted on the indexed attribute, making it the practical choice for very large, sorted data files where index memory footprint matters.

Q8. Indexed Sequential File Organization vs Direct File Organization

Bloom's Level: 4 – Analyze

Comparison Table

Aspect	Indexed Sequential File Organization	Direct (Random) File Organization
Access method	Combination of sequential access (data sorted by key) plus an index for direct access	Records accessed directly via a computed address (usually a hash function on the key)
Data ordering	Physically sorted on the key field	No particular physical order required
Retrieval	Supports both sequential (range) scans and fast direct lookup via index	Very fast direct/point lookup, but poor for sequential/range access
Insertion/Deletion	Moderate cost — may need overflow areas and periodic reorganization to keep the index and order consistent	Generally fast, since address is computed directly, but collisions may need handling

Aspect	Indexed Sequential File Organization	Direct (Random) File Organization
Storage overhead	Extra space needed for the index structure	Minimal extra structure, though hash collisions may waste space
Typical use case	Banking, payroll — need both batch (sequential) processing and individual account lookup	Real-time transaction systems needing fast single-record access, e.g., session lookup

Explanation

Indexed sequential file organization keeps the data file physically sorted by key (like sequential organization) but adds an index on top, giving the best of both worlds: efficient sequential/batch processing and reasonably fast direct access via the index, at the cost of extra index storage and reorganization overhead when the file changes frequently.

Direct (random / hashed) file organization computes a record's storage address directly from its key using a hash function, giving very fast point lookups without needing an index, but it sacrifices any notion of physical ordering, making range queries and sequential processing inefficient.

Example(s)

- A bank's account master file is often indexed-sequential — sorted by account number for month-end batch reports, but indexed for a teller to instantly pull up one account.
- A session-store keyed by a randomly generated session ID benefits from direct/hashed organization since sessions are always looked up individually, never scanned as a range.

Justification / Critical Analysis

Indexed sequential organization is preferable when an application needs a mix of range/batch processing and individual lookups, whereas direct file organization is preferable when the workload is dominated by single-key point access and ordering/range queries are unimportant. The limitation of indexed-sequential is the overhead of maintaining both the sort order and the index during frequent updates; the limitation of direct organization is its complete unsuitability for range-based or sorted reporting.

Q9. Single-Level Indexing vs Multi-Level Indexing

Bloom's Level: 4 – Analyze

Comparison Table

Aspect	Single-Level Indexing	Multi-Level Indexing
Structure	One flat index mapping key values directly to data block/record addresses	A hierarchy of indexes — an index built on top of another index (e.g., outer index, inner index)
Search cost	Search time grows with index size ($O(\log n)$ with binary search, but the whole index may not fit in memory)	Search cost reduced further since only the small top-level index needs scanning to locate the relevant portion of the lower index
Suitability for large files	Becomes inefficient as the file — and hence the index — grows very large	Scales well for very large files since higher levels stay small enough to fit in memory

Aspect	Single-Level Indexing	Multi-Level Indexing
Structure basis	Simple linear or binary-searchable list	Typically implemented as a tree structure (e.g., B+ Tree), often multiple levels deep
Maintenance complexity	Simpler to build and maintain	More complex — insertions/deletions may cascade changes across levels

Explanation

A single-level index is one flat structure directly mapping key values to record/block addresses. While a binary search on this index is efficient in theory, once the file becomes very large, even the index itself may become too large to fit in memory, slowing down every access.

Multi-level indexing addresses this by building an index on the index — a smaller, higher-level (outer) index that helps quickly locate the relevant section of the lower-level (inner) index. This hierarchy can continue for as many levels as needed, and is the principle underlying B-Tree/B+ Tree indexes used in modern databases.

Example(s)

- A single-level index on a modest 10,000-row table performs well since the index itself is small.
- A multi-level (B+ Tree) index on a table with 100 million rows keeps the tree only 3–4 levels deep, so any record can be found in just 3–4 disk I/Os instead of a large binary search over a huge flat index.

Justification / Critical Analysis

As data volume grows, a single-level index eventually loses its efficiency advantage because the index itself becomes too large for fast memory-resident search. Multi-level indexing directly solves this scalability problem by keeping the top levels of the hierarchy small and memory-resident, which is precisely why virtually all production-grade relational databases use multi-level tree-based indexes (B+ Trees) rather than flat single-level indexes for large tables.

Q10. Hashing Techniques vs Indexing Techniques for Exact-Match and Range Queries

Bloom's Level: 5 – Evaluate

Comparison Table

Aspect	Hashing Techniques	Indexing Techniques (B-Tree / B+ Tree)
Exact-match queries	Excellent — average $O(1)$ lookup by computing the hash directly	Good — $O(\log n)$, slower than hashing but still efficient
Range queries (e.g., BETWEEN, >, <)	Poor — hash values do not preserve order, so ranges require scanning many/all buckets	Excellent — sorted tree/leaf structure allows efficient sequential range scans
Sorting / ORDER BY support	Not supported directly — no inherent ordering	Naturally supported since data/keys are kept sorted
Worst-case performance	Degrades under heavy collisions/overflow chains to $O(n)$	Bounded — $O(\log n)$ even in the worst case for balanced trees

Aspect	Hashing Techniques	Indexing Techniques (B-Tree / B+ Tree)
Best-fit scenario	Point lookups on high-cardinality unique keys (e.g., primary key equality search)	Range-heavy, analytical, and reporting workloads

Explanation

For exact-match queries, hashing techniques are generally superior because a well-designed hash function computes the target bucket directly in constant average time, avoiding the multi-level traversal that a tree index requires.

For range queries, indexing techniques (particularly B+ Trees) clearly outperform hashing, because hashing intentionally scrambles key order to distribute values evenly across buckets — this is exactly what makes it fast for equality lookups but useless for retrieving all values between two bounds, since matching records could be scattered across every bucket.

Example(s)

- Evidence from the SQLite demonstration (see Q2) confirms this: an indexed range query on `roll_no` (`SEARCH ... rowid>? AND rowid<?`) used the ordered B-Tree-like primary key structure efficiently, something a pure hash index cannot do without a full bucket scan.
- A database's `WHERE customer_id = 12345` (equality) is ideal for a hash index, whereas `WHERE order_date BETWEEN '2026-01-01' AND '2026-06-30'` (range) requires a B+ Tree index to be efficient.

Justification / Critical Analysis

Evaluating both techniques together: neither is universally superior — the correct choice depends entirely on the query workload. Hashing should be chosen when the dominant access pattern is exact-match lookups on high-cardinality keys, since it offers near-constant-time performance unmatched by tree indexes. Indexing (B-Tree/B+ Tree) should be chosen whenever range queries, sorting, or partial-match queries are common, since only an ordered structure supports these efficiently. Many production databases pragmatically use both: a B+ Tree as the primary/clustered structure for ordering and range support, supplemented with hash indexes on specific high-traffic equality-lookup columns.